

Soft vs Rigid Gripper Benchmark for Cloth Grasping

A Force-Depth-Tilt Characterization of Parallel-Jaw and Finray Geometries

Shiyao Ni
Master of Science



Department of Mechanical Engineering
McGill University
Montreal, Quebec, Canada

June 2026

A thesis submitted to McGill University in partial
fulfillment of the requirements of the degree of
Master of Science

© Shiyao Ni, 2026

Abstract

Robotic grasping of thin, deformable fabric remains an unsolved problem in industrial automation. Two competing gripper geometries are routinely proposed: rigid parallel-jaw fingers, which offer precise force control but a narrow operational tolerance, and compliant finray-effect fingers, which trade force precision for a wider operational envelope. This thesis presents an experimental benchmark of six gripper variants—two parallel-jaw configurations (rigid stock, and stock plus a Shore-90A thermoplastic polyurethane pad) and four finray configurations (18 mm and 26 mm finger thicknesses, a 5° inclined-seat fixture, and a second print of the 18 mm finger)—mounted on a Franka FR3 collaborative manipulator. Each variant was characterised across a two-dimensional input space of press depth (the distance the gripper fingertip is driven below the detected cloth surface) and tilt angle (the gripper orientation about world x and y axes relative to the surface normal). Stationary contact force in the gripper z -axis was sampled during each grasp and used as the primary outcome variable, with binary grasp success as a secondary outcome.

The benchmark produced 845 analysed trials across the six variants. The rigid parallel-jaw variant exhibits a narrow 1.5 mm working depth range bounded above by a -60 N safety reflex and below by no-contact; the TPU-padded variant widens this to 2.5 mm. The finray variants, by contrast, remain successful across the full 10 mm calibrated analysis window (working range ≥ 10 mm) with a soft force plateau past the buckling knee (peak forces of 30–51 N across the finray variants).

Beyond this expected compliance advantage, two depth-dependent tilt effects emerged. First, finray contact force is strongly tilt-sensitive at the no-contact boundary but robust to tilt in the soft plateau: a $\pm 5^\circ$ tilt produces a 58% contact force reduction at -0.5 mm depth but only 1% at $+10$ mm depth. Second, the finray grip-

per exhibits a direction-dependent handedness under tilt about its asymmetric (y) axis at boundary depths—a $\pm 2.5^\circ$ inversion reverses the contact force by a factor of $2.36\times$ at -1.0 mm —but this handedness vanishes at safe operating depths. We propose this as an emergent characterisation of the finray geometry: the deeply pressed finger structure wraps the contact surface symmetrically regardless of tilt sign, while at lighter contact the asymmetric ridge profile dominates the force response.

The implication for cloth-handling system design is that compliant grippers achieve directional robustness only when pressed past a depth threshold; control policies that operate near the no-contact boundary for delicate fabric handling must explicitly model the gripper’s directional asymmetry.

Abrégé

La préhension robotique de tissus minces et déformables demeure un problème non résolu en automatisation industrielle. Deux géométries de préhenseur concurrentes sont couramment proposées : les doigts rigides à mâchoires parallèles, qui offrent un contrôle de force précis mais une tolérance opérationnelle étroite, et les doigts souples à effet finray, qui échangent la précision de force contre une enveloppe opérationnelle plus large. Cette thèse présente un banc d'essai expérimental de six variantes de préhenseur — deux configurations à mâchoires parallèles (rigide d'origine, et d'origine avec un coussinet en polyuréthane thermoplastique de dureté Shore 90A) et quatre configurations finray (doigts de 18 mm et de 26 mm d'épaisseur, une fixation inclinée à 5° , et une seconde impression du doigt de 18 mm) — montées sur un manipulateur collaboratif Franka FR3. Chaque variante a été caractérisée sur un espace d'entrée bidimensionnel composé de la profondeur d'appui (la distance à laquelle le bout du doigt est enfoncé sous la surface détectée du tissu) et de l'angle d'inclinaison (l'orientation du préhenseur autour des axes x et y du repère monde par rapport à la normale de la surface). La force de contact stationnaire selon l'axe z du préhenseur est échantillonnée pendant chaque prise et sert de variable de résultat principale, la réussite binaire de la prise servant de résultat secondaire.

Le banc d'essai a produit 845 essais analysés sur les six variantes. Le préhenseur rigide à mâchoires parallèles présente une plage de profondeur utile étroite de 1,5 mm, bornée supérieurement par un réflexe de sécurité à -60 N et inférieurement par l'absence de contact ; la variante avec coussinet TPU élargit cette plage à 2,5 mm. Les variantes finray, en revanche, restent efficaces sur toute la fenêtre d'analyse calibrée de 10 mm (plage utile de 10 mm ou plus), avec un plateau de force souple au-delà du pic de flambement (forces maximales de 30 à 51 N selon la variante).

Au-delà de cet avantage attendu de la souplesse, deux effets d’inclinaison dépendants de la profondeur ont émergé. Premièrement, la force de contact du finray est fortement sensible à l’inclinaison près de la limite de non-contact, mais robuste à l’inclinaison dans le plateau souple : une inclinaison de $\pm 5^\circ$ produit une réduction de 58 % de la force de contact à une profondeur de $-0,5$ mm, mais de seulement 1 % à $+10$ mm. Deuxièmement, le préhenseur finray présente une réponse directionnelle (“chiralité”) sous inclinaison autour de son axe asymétrique (y) aux profondeurs limites — une inversion de $\pm 2,5^\circ$ inverse la force de contact d’un facteur de 2,36 à $-1,0$ mm — mais cette chiralité disparaît aux profondeurs d’opération sûres. Nous proposons ceci comme une caractérisation émergente de la géométrie finray : le doigt profondément appuyé épouse la surface de contact de manière symétrique quel que soit le signe de l’inclinaison, tandis qu’à contact plus léger, le profil asymétrique de l’arête domine la réponse en force.

L’implication pour la conception de systèmes de manipulation de tissus est que les préhenseurs souples n’atteignent une robustesse directionnelle qu’au-delà d’un seuil de profondeur d’appui ; les politiques de commande qui opèrent près de la limite de non-contact pour la manipulation délicate de tissus doivent explicitement modéliser l’asymétrie directionnelle du préhenseur.

[TODO: Faire vérifier cette traduction par un collègue francophone avant le dépôt final (exigence de qualité, non de conformité — GPS exige seulement la présence d’un abrégé en français).]

Contribution to Original Knowledge

This thesis contributes the first controlled, parametric characterisation of how compliant finray-effect fingers and rigid parallel-jaw grippers grasp cloth as a function of press depth and gripper tilt, holding the robot, control policy, cloth specimen, and environment fixed and varying only the end-effector. Prior work on finray-family fingers has modelled or measured *lateral* grasping contact; no published model or benchmark, to the author’s knowledge, characterises the axial tip-compression response — the rise–peak–valley–recovery force profile identified here, shared across four finray variants to within 1% normalised error. The principal empirical contribution is the demonstration that finray tilt robustness is *depth-gated*: strongly tilt-sensitive near the no-contact boundary, but robust to tilt once pressed past the buckling threshold, with a direction-dependent asymmetry about the finger’s geometrically asymmetric axis that likewise vanishes in the safe operating regime. This depth-dependent tilt robustness had not been quantitatively characterised in prior literature. Secondary contributions include a reproducible benchmark protocol and its open-source implementation, a documented per-variant tool-centre-point calibration procedure, and a geometric verification framework disentangling control-frame positioning error from physical gripper compliance under load.

Contribution of Authors

This thesis is the original work of the author, Shiyao Ni, under the supervision of Prof. Audrey Sedal. The author carried out all experimental design, hardware integration, software development, data collection, analysis, and writing.

The Franka FR3 manipulator and its ROS 2 control stack were maintained by the MACRObotics Lab. The 3D-printed finray fingers were fabricated using lab resources at the MACRObotics Lab. No portion of this thesis has been previously published.

This is a traditional (monograph-style) thesis; no chapter derives from a co-authored manuscript, and no collaborator contributed content requiring declaration under McGill GPS guidelines beyond the infrastructure acknowledgements above.

Acknowledgements

I thank my supervisor, Prof. Audrey Sedal, for her guidance throughout this work, and the members of the MACRObotics Lab at McGill University for their support in the laboratory. I am grateful to the McGill Applied Dynamics Lab for developing and maintaining the `fr3_pose_controller` ROS2 control stack on which every experiment in this thesis depends, and to the broader ROS2 and Franka robotics communities whose open-source tooling made this work possible.

Finally, I thank my family for their patience and encouragement.

[TODO: Add funding sources (scholarships, grants, or stipend support), if any, per McGill GPS acknowledgement conventions.]

Contents

Abstract	i
Abrégé	iii
Contribution to Original Knowledge	v
Contribution of Authors	vi
Acknowledgements	vii
List of Abbreviations	xiii
Notation	xiv
Preface	xvii
1 Introduction	1
1.1 Motivation	1
1.2 Research Questions and Contributions	2
1.3 Thesis Organisation	3
2 Background and Related Work	4
2.1 Cloth Grasping in Robotics	4
2.2 Strategies for Acquiring the First Grasp	5
2.2.1 Environment-Assisted and State-Changing Actions	6
2.2.2 Specialised End-Effectors	6
2.2.3 The Gap: the Gripper as an Experimental Variable	7

2.3	Compliant Gripper Designs	8
2.4	Force-Controlled Grasping	10
2.5	The Franka FR3 Platform	11
2.6	Position and Summary	12
3	Experimental Setup and Methodology	14
3.1	Hardware Configuration	14
3.1.1	Robot Manipulator	14
3.1.2	Gripper Variants Tested	15
3.1.3	Cloth Specimen	16
3.1.4	Test Surface	16
3.2	Reference Frame and Tool-Centre Calibration	16
3.2.1	Frame Convention	16
3.2.2	TCP Offset Calibration	17
3.3	Force Measurement	18
3.4	Tilt-Aware Pose Targeting	18
3.5	Data Pipeline	19
3.5.1	Per-Trial Recording	19
3.5.2	Per-Session Recording	19
3.5.3	Verdict Recording and Outcome Classification	20
3.6	Trial Procedure	20
3.7	Statistical Treatment	21
3.7.1	Calibrated Depth Reference	21
3.7.2	Per-Cell Statistics	22
4	Experimental Design and Procedure	27
4.1	Study Design	27
4.2	Depth Grid	28
4.3	Tilt-Aware Pose Geometry: Verification	29
4.4	Rotation Tracking Verification	29
4.5	Trial Protocol and Outcome Classification	30
4.6	Experiment 3: Inclined-Surface Demonstration (Pilot)	31

5	Results	32
5.1	Canonical Force-Depth Curves at 0°	32
5.2	Bipolar Tilt Bias on the Finray Family	34
5.3	Cross-Axis Sign Reversal at the Boundary	35
5.4	Mount-Tilt Alignment: 5° Mount as a Depth Shift	35
5.5	Tilt-Aligned Response Comparison	36
5.6	Print-to-Print Manufacturing Variability	36
5.7	Trial Outcome Distribution	37
5.8	Working-Range and Failure Comparison	38
6	Discussion	41
6.1	Bell-Plus-Valley Curve Shape	41
6.2	Cross-Axis Sign Reversal: Mechanism	42
6.3	Mount-Tilt Equivalence: Geometric Interpretation	43
6.4	Manufacturing Variability vs Session Drift	43
6.5	Failure-Mode Taxonomy: Why the Distinction Matters	44
6.6	Deployment-Envelope Implications for Cloth-Handling Systems	45
6.7	Limitations	45
6.8	Code and Data Availability	46
7	Conclusion and Future Work	47
7.1	Answers to the Research Questions	47
7.2	Implications	48
7.3	Future Work	49
A	Detailed Coverage and Supplementary Data	51
A.1	Per-Variant Coverage Table	51
	List of Publications	54

List of Figures

3.1	Gripper drawings and design parameters α and t	23
3.2	Cloth specimen	24
3.3	Runtime tilt convention	25
3.4	Finray engagement at press depth	26
5.1	Shared force-depth shape across the finray family	34
5.2	Tilt-aligned response per finray variant	36
5.3	Print-to-print force comparison	37
5.4	Grasp success rate versus calibrated depth	38
5.5	Zero-tilt working range per variant	39
5.6	Working range under tilt, all variants	40

List of Tables

3.1	Gripper variants	15
3.2	Calibrated TCP offsets	17
4.1	Test conditions	28
4.2	Tilt geometry verification	29
5.1	Zero-tilt curve summary	33
5.2	Pairwise shape agreement	33
5.3	Cross-axis bias sign reversal at boundary	35
5.4	Mount tilt equivalence	35
5.5	No-grasp fractions at boundary	38
A.1	Experiment 1 coverage	52
A.2	Experiment 2 coverage	53

List of Abbreviations

API	Application Programming Interface
BOM	Bill of Materials
CAD	Computer-Aided Design
DOF	Degree(s) of Freedom
FCI	Franka Control Interface
FR3	Franka Research 3 (robot model)
FSM	Finite-State Machine
GT	Ground Truth (calibrated surface height record)
PEI	Polyetherimide (3D-printer build-plate surface)
RGB-D	Red–Green–Blue plus Depth (camera modality)
ROS 2	Robot Operating System 2
RPY	Roll–Pitch–Yaw (orientation angles)
TCP	Tool Centre Point
TPU	Thermoplastic Polyurethane
URDF	Unified Robot Description Format

Notation

Frames and poses

$\mathcal{W}, \mathcal{F}, \mathcal{T}$	World (robot base), gripper flange, and gripper tip frames
$\mathbf{p}$	Position vector in $\mathbb{R}^3$ (boldface); $\mathbf{p}_{\mathcal{F}}, \mathbf{p}_{\mathcal{T}}$ denote the flange and tip positions expressed in $\mathcal{W}$
R, R_0	Orientation of the gripper as a rotation matrix, $R \in SO(3)$; $R_0 = R_x(180^\circ)$ is the nominal down-pointing orientation
$R_x(\phi), R_y(\phi)$	Elementary rotation matrix for a rotation by angle ϕ about the fixed world x or y axis
$\hat{\mathbf{z}}$	Unit vector along $+z$
z_{TCP}	Per-variant tool-centre-point offset from flange to tip along the flange's local z axis (mm); $z_{\text{TCP}}^{(v)}$ denotes the calibrated value for variant v
x_0, y_0	Fixed probe/target position in the world xy plane (mm)

Design parameters (fixed per variant)

t	Finray finger blade thickness (mm) (Fig. 3.1)
α	Baseline seat-mount tilt angle, about the gripper's local y axis, introduced by the 3D-printed inclined seat ($^\circ$) (Fig. 3.1)

Runtime and trial variables

θ_x, θ_y	Commanded robot-arm tilt angles about the world x / y axes (which coincide with the gripper's local x / y axis lines at R_0), applied as $R_{\text{tilt}} = R_y(\theta_y) R_x(\theta_x)$ (extrinsic) ($^\circ$); distinct from, and never summed with, the seat angle α (Fig. 3.3)
d	Press depth: distance the gripper fingertip is driven below the detected cloth/table surface (mm)
d_{cal}	Calibrated press depth: $d - d_{\text{first success}}$, zeroed at each series' first majority-success cell (mm)
z_{surf}	Detected cloth/table surface height at the tip's xy location (mm)
a	Tilt axis identifier (x or y)
$b(d, a)$	Signed tilt bias: $b(d, a) = F_z^{+2.5^\circ}(d, a) - F_z^{-2.5^\circ}(d, a) $ (N)

Force measurement

$\mathbf{F}_{\text{ext}}$	Franka external wrench force vector at the end-effector flange (N); F_x, F_y, F_z denote its components
F_z^{bias}	Free-air F_z bias captured at session start (N)
$F_z^{(i)}$	The i -th stationary F_z sample within a trial (N)
N	Number of stationary force samples pooled in a (variant, depth, tilt) cell
$\bar{F}_z, \sigma_{F_z}$	Per-cell mean and standard deviation of F_z over pooled trials
$\sigma_{z,\text{tcp}}$	Across-tap standard deviation of the touch-off z measurement during TCP calibration (mm; distinct from σ_{F_z} above)

Derived quantities

L_{eff}	Effective finger lever-arm length inferred from the mount-tilt/depth-shift equivalence (mm)
------------------	--

Preface

This work grew out of a practical frustration: in the laboratory, as in much of the published literature, robotic cloth-manipulation demonstrations routinely begin with a human placing the fabric into the robot’s gripper. The step that most limits autonomy — acquiring the first grasp of a thin, flat cloth lying on a surface — is the step most often skipped. Rather than approaching that gap with more perception or learning, this thesis asks a hardware question: how much of the problem dissolves if the gripper itself is compliant, and what, precisely, is traded away in exchange?

All chapters of this thesis are original, previously unpublished work by the author. Collaborative and infrastructural contributions are declared in the Contribution of Authors section.

1

Introduction

1.1 Motivation

Fabric handling automation is a long-standing open problem in industrial robotics, with applications spanning apparel manufacturing and textile logistics, laundry processing and folding, hospital bedding management, assistive robotics for elderly or disabled individuals, and the packaging of soft goods. Yet a striking fraction of published robot-manipulation demonstrations sidesteps the hardest step entirely: the object is handed to the robot’s gripper by a human operator, and the demonstrated autonomy begins only after the grasp already exists.¹ For cloth specifically, acquiring the first grasp (picking a thin, flat fabric directly off a surface) remains the bottleneck that limits how autonomous these systems can truly be. Cloth is uniquely difficult to grasp robustly: it deforms continuously under contact, exhibits no fixed graspable feature, and has a near-zero bending stiffness that allows it to slip out of a poorly

¹A prominent recent example: a widely circulated January 2024 video of the Tesla Optimus humanoid folding a shirt was subsequently acknowledged by its poster to be teleoperated, with the shirt pre-staged in a fixed setup [1, 2].

placed gripper. Conventional rigid-jaw grippers, designed for sturdy objects with well-defined geometry, perform poorly on cloth without elaborate tactile sensing or human-tuned approach trajectories.

A class of compliant grippers based on the finray effect, a passive bending mechanism inspired by fish-fin anatomy, has been proposed as a solution. Finray fingers passively wrap around the contact surface, increasing contact area and gripping force without requiring active force control. This compliance, however, comes at the cost of force-modelling difficulty: the finger geometry deforms under load, the contact point shifts, and tilt-induced asymmetries arise that have no analogue in the rigid-jaw case.

The central question of this thesis is: *to what extent does the finray’s compliance trade single-axis force precision for multi-axis operational robustness?* Specifically, how does the finray’s contact force respond to deviations in the two design degrees of freedom most likely to be uncontrolled in real deployment: press depth and gripper tilt angle?

1.2 Research Questions and Contributions

This thesis answers the following research questions through a controlled experimental benchmark on a Franka FR3 collaborative robot. The benchmark treats the gripper itself as the experimental variable, holding the robot, control policy, cloth, and environment identical across all conditions, and asks:

1. How does the contact force vary with press depth, for each finger design?
2. How much can the gripper be tilted before the grasp becomes unreliable?
3. Does tilt tolerance depend on the press depth?
4. What is the soft gripper’s response under compression?

The contributions of this work are:

- A reproducible benchmark protocol for cloth-grasp force characterisation on a Franka FR3, with full source code published alongside this thesis.

- A two-dimensional depth–tilt experimental matrix sampling six gripper variants: two rigid parallel-jaw configurations and four finray prints spanning blade thickness, seat angle, and print instance.
- Empirical evidence of a depth-dependent tilt robustness transition in finray grippers, with quantitative characterisation of the transition depth.
- A documented calibration procedure for per-variant tool-centre-point (TCP) offset using differential force touch-off against a rigid reference.
- A geometric verification framework that disentangles control-frame positioning error from physical gripper compliance under load.

Scope. The benchmark deliberately isolates the *contact* phase of cloth grasping, the interval between first gripper–cloth contact and a completed lift, and characterises it over a continuous depth–tilt parameter space. Downstream task-level evaluation (folding, placing, and dragging under the same policy) is future work (Chapter 7). This narrowing is intentional: it is precisely because the contact phase is usually bundled inside end-to-end task pipelines that its mechanics have not previously been mapped as a function of gripper design.

1.3 Thesis Organisation

Chapter 2 reviews the literature on cloth grasping, compliant grippers, and force-controlled approach strategies. Chapter 3 describes the experimental hardware, the kinematic calibration procedure, and the data-acquisition pipeline. Chapter 4 presents the two experimental campaigns: the depth-sweep characterisation (Experiment 1) across all six gripper variants at zero tilt, and the tilt-sweep characterisation (Experiment 2) at boundary and safe depths. Chapter 5 reports the resulting force–depth–tilt characterisation. Chapter 6 synthesises the findings into a model of depth-dependent tilt robustness. Chapter 7 concludes and identifies directions for future work.

Background and Related Work

2.1 Cloth Grasping in Robotics

Robotic manipulation of cloth and other deformable fabric is difficult precisely because the object being manipulated lacks a fixed geometry: unlike a rigid part, a cloth's shape is itself part of the manipulation state, and it must typically be inferred from vision rather than assumed from a physical model. Early work in this space focused on perceiving that shape well enough to act on it. Kita et al. [3] fit a deformable mass-spring model of a held garment to stereo-camera silhouette observations, using the best-matching simulated deformation to estimate the position of an unheld part of the garment for a second manipulator to grasp. This was an early demonstration that a physics-informed shape prior, rather than raw image features, is what makes cloth state estimation tractable. Maitin-Shepard et al. [4] took a complementary, model-light approach: a multiple-view geometric-cue detector locates cloth corners directly from silhouettes, in a manner robust to texture and colour variation, letting a general-purpose PR2 manipulator pick and fold towels of varying material and ap-

pearance without a specialised end-effector. Miller et al. [5] generalised this line of work with a polygonal-approximation geometric model that tracks a garment’s silhouette through an entire folding sequence, allowing one parameterised folding strategy to transfer across previously unseen towels, pants, and shirts. Li et al. [6] pushed the physical modelling further still, using a predictive thin-shell simulation fit to RGB-D observations to plan regrasp actions that unfold and flatten a crumpled garment by selecting the grasp points expected to maximise resulting flatness.

More recent work has moved from hand-designed geometric priors toward learned policies and richer sensing. Saxena and Shibata [7] trained convolutional classifiers to recognise garment type and locate hem/waistline grasp points directly from images, as the perception front-end of a robotic dressing-assistance pipeline. Wu et al. [8] showed that a model-free visual reinforcement-learning policy for cloth and rope manipulation can be learned an order of magnitude faster than a jointly-learned pick-and-place policy by conditioning only on the *placing* action and selecting picks that maximise the value of the resulting placement, entirely without human demonstrations. Closest in spirit to the force-sensing emphasis of this thesis, Tirumala et al. [9] mounted a tactile sensor on a Franka gripper fingertip and trained a classifier on the resulting tactile signal to determine how many layers of cloth had been grasped, using that signal to re-grasp until exactly one or two layers were singulated from a stack. This is direct evidence that certain cloth-grasp ambiguities cannot be resolved by vision alone. For a comprehensive treatment of the field beyond these representative examples, Longhini et al. [10] provide a recent survey spanning perception, dynamics modelling, benchmarking, and control specifically for textile manipulation.

2.2 Strategies for Acquiring the First Grasp

Because a flat cloth on a surface presents no protruding feature for a conventional pinch grasp, the literature has developed two broad families of workarounds: manipulating the *environment or cloth state* until a graspable feature exists, and replacing the *end-effector* with a mechanism that does not require one.

2.2.1 Environment-Assisted and State-Changing Actions

The first family keeps a standard gripper and instead spends actions creating a graspable feature (a fold, wrinkle, or free edge) before the pick. Ramisa et al. [11] detected and ranked wrinkles in depth images of crumpled laundry, directing a standard gripper to the most graspable bump on the cloth rather than to a flat region. Qian et al. [12] segmented cloth edges and corners (as distinct from wrinkles and folds) in depth images and coupled the segmentation to a sliding grasp policy that approaches along the estimated grasp direction. At the dynamic extreme, Ha and Song’s Fling-Bot [13] showed that a high-velocity fling primitive can unfold a crumpled cloth in a few actions: environment interaction used not to create a pinchable feature but to reset the cloth to a flatter, more predictable state. These strategies are effective, but each spends additional actions and sensing before the grasp itself, and their reliability degrades exactly in the flat-cloth-on-surface configuration this thesis targets, where there is no wrinkle to find and no free edge to slide against.

2.2.2 Specialised End-Effectors

The second family redesigns the gripper so that a protruding feature is unnecessary. Monkman’s survey of grippers for fibrous materials [14] catalogues the classical industrial approaches, developed largely for garment manufacturing: intrusive (needle/pin) grippers that pierce the top ply, and adhesive or electrostatic surface grippers. Electrode adhesion has since matured substantially; Guo et al. [15] review its use across robotics and highlight its suitability for delicate, flexible materials at low energy cost. Non-contact airflow methods lift fabric without touching it at all: Ozcelik and Erzincanli’s Bernoulli-principle end-effector [16] suspends a garment ply on a radial airflow film. More recently, variable-friction fingertips blur the line between gripper and manipulation policy: Jiang et al. [17] switch fingertip surfaces between high- and low-friction states to slide, pinch, and unfold fabric in-hand. These end-effectors trade generality for capability: needles risk damaging fine textiles, electrode adhesion and airflow devices require dedicated power/pneumatic infrastructure and are largely single-purpose, and all of them abandon the ubiquitous two-jaw form factor that general-purpose robot platforms (including the Franka Hand used here) provide.

2.2.3 The Gap: the Gripper as an Experimental Variable

Across both families, the gripper hardware is treated as *fixed* and the innovation budget is spent on perception, policy, or a task-specific mechanism. Direct comparisons between compliant and rigid jaw-type grippers exist; Li et al. [18] systematically review soft-rigid hybrid grippers against industrial parallel-jaw grippers across payload, adaptability, and precision metrics. Such comparisons, however, are conducted on rigid or quasi-rigid workpieces, not on cloth, and report capability envelopes rather than parametric contact behaviour.

The closest precedent to this thesis is the household-clothing benchmark of Clark et al. [19], who evaluated several commercial end-effectors (including finray-style compliant fingers mounted on a Franka Emika Hand, alongside the standard Franka Hand, a Robotiq 2F-140, and the OpenHand Model T42) on four cloth-specific benchmarks: edge-drag accuracy, edge-grasp resilience, crumpled-object encapsulation, and fold generation. Their results establish empirically that end-effector selection materially changes cloth-manipulation outcomes, and their inclusion of finray fingers on the same robot platform used here makes the comparison direct. The evaluation, however, is scored at the *task* level: each benchmark reports success or accuracy of a completed manoeuvre, with the approach pose fixed per task. What it does not measure is how the grasp itself degrades as the approach conditions vary: the continuous dependence of contact force and grasp outcome on press depth and gripper tilt that any deployed system experiences as positioning error.

None of the work above, therefore, addresses the specific question this thesis investigates: cloth-manipulation literature treats grasp *success* as an outcome to be predicted or achieved at a fixed operating point, not as a function of a continuous physical parameter space (press depth, gripper tilt) whose boundary can be mapped and compared across gripper geometries. The benchmark in this thesis is deliberately narrower in manipulation scope than folding or dressing pipelines, but correspondingly more quantitative about the mechanics of the grasp itself: it keeps the two-jaw form factor fixed, varies only the finger design (t, α) (defined in Figure 3.1), and maps the resulting contact-force landscape under the positional and angular uncertainty (d, θ_x, θ_y) that any deployed system must tolerate.

2.3 Compliant Gripper Designs

The finray-effect finger design used in this thesis originates in a patent by Kniese [20], who observed that a structural element built from two flexible longitudinal members joined by hinged cross-struts deflects *toward* an applied lateral load rather than away from it, the counter-intuitive “bends toward the pressure” behaviour later commercialised by Festo as the Fin Ray Effect®. Festo’s own current technical documentation for its DHAS gripper-finger product line [21] describes this structure explicitly as two flexible bands joined by cross-ribs at flex hinges, and reports measured retention-force, lateral-force, and indentation-depth curves as functions of gripping force and workpiece diameter. These are the same three quantities (force, indentation depth, contact geometry) that this thesis’s depth-sweep experiments characterise directly on 3D-printed finray fingers rather than the commercial product.

Academic characterisation of the finray design has followed two directions: analytic modelling and empirical/manufacturing studies. Shan and Birglen [22] developed an analytic mechanical model of finray fingers to predict grasp force and finger deformation from geometric parameters, addressing the absence of a convenient design tool for this class of compliant finger. Their model, like the empirical studies that follow it, addresses *lateral* grasping contact; no published model, to the author’s knowledge, predicts the axial tip-compression response (the rise–peak–valley–recovery profile) characterised empirically in this thesis. On the empirical side, Crooks et al. [23] modified the finray design by angling its internal crossbeams to introduce a directional bending bias, fabricated by multi-material 3D printing, and reported an approximately 40% increase in holding capacity over a traditional finray finger from this geometric change alone. Elgeneidy et al. [24] characterised fully 3D-printed finray fingers made from flexible thermoplastic polyurethane (TPU) filament, varying cross-beam rib angle and spacing, and found that at large tip displacements the ribs contact one another (*layer jamming*), substantially increasing force generation and producing a tunable-stiffness effect. Most directly relevant to this thesis’s use of TPU-printed fingers, Lang et al. [25] 3D-printed finray fingers in both TPU 60A and TPU 95A and tensile-tested them: the softer 60A durometer gave the better force/compliance balance and left delicate fruit undamaged, while 95A reached substantially higher forces

but visibly bruised the fruit. That result is direct prior evidence that finger material hardness, independent of finger geometry, is a first-order design lever for a compliant gripper, a finding this thesis’s own manufacturing-variability results (Section 5.6) are consistent with.

The operational-envelope argument for compliant hands (that compliance buys tolerance to positioning error) was established quantitatively by Dollar and Howe’s SDM Hand [26]: an underactuated hand with compliant polymer fingers that grasped a 48 mm cylinder at positions offset by more than 100% of the object size in each direction, while keeping acquisition contact forces low, where a rigid gripper would require millimetre-level placement. Their successful-grasp-range maps over an xy offset grid are the direct methodological ancestor of the depth–tilt working-window maps produced in this thesis, transposed from rigid household objects to cloth and from lateral offset to press depth and tilt. Closest of all in spirit, Teeple et al. [27] decomposed the benefit of a pneumatic soft gripper’s compliance into per-axis contributions when handling thin, flexible materials: vertical compliance was shown to buy robustness to large vertical positioning uncertainty, measured head-to-head against a rigid gripper, while lateral and rotational compliance lengthen the response window and cap the tensile force during snags. Their vertical-uncertainty comparison is the single-gripper-pair analogue of this thesis’s central question; the benchmark here extends it across a family of finger designs, adds the tilt dimension, and grounds it on cloth with a fully passive compliant finger.

Compliant gripping is not achieved solely through finray-style flexure. Granular jamming is the other major compliant-gripping paradigm in the literature: Brown et al. [28] demonstrated that a latex membrane filled with granular material conforms to an object’s shape under positive pressure and rigidifies under vacuum, enabling passive universal grasping without dedicated fingers at all, and Amend et al. [29] extended this with a combined positive/negative pressure actuation cycle for faster, more reliable grasp-and-release cycles. A third paradigm, pneumatic elastomer actuation, underlies commercial soft grippers such as those from Soft Robotics Inc.; the technology traces to pneumatic-network (PneuNet) elastomer actuators developed by Ilievski et al. [30], and the commercial mGrip product line [31] demonstrates that compliant gripping has moved beyond academic prototypes into food-handling

industrial deployment. This thesis’s finray variants and rigid parallel-jaw comparator sit within this broader design space as one specific, geometrically simple point in it, chosen because (unlike jamming or pneumatic actuation) a finray finger requires no additional actuation hardware beyond the gripper’s own jaw motion.

2.4 Force-Controlled Grasping

The theoretical basis for treating contact force, rather than position alone, as a control variable during manipulation is Hogan’s impedance control framework [32, 33, 34]: neither pure position control nor pure force control is adequate for a task involving contact with an environment of unknown or varying mechanical properties, so the manipulator should instead regulate the dynamic relationship — the impedance — between its motion and the interaction force at its end-point. This thesis’s trial procedure, which presses the gripper to a commanded depth and then reports the resulting *stationary* contact force F_z rather than commanding force directly, is an impedance-control-style measurement: depth is the commanded quantity, force is the observed response, and the relationship between them (Chapter 5) is exactly the impedance curve Hogan’s framework predicts should characterise contact with an unknown environment (here, the cloth-covered rigid table).

On the Franka FR3 platform specifically, libfranka exposes two real-time command modes: external torque-level control (including a Cartesian impedance controller distributed as a reference implementation) and Cartesian pose motion generation, in which the client streams pose targets at the 1 kHz servo rate and tracking is delegated to the arm’s internal impedance controller [35]. The tilt-aware pose targeting used in this thesis’s trial procedure (Section 3.4) uses the latter mode: the in-house `fr3_pose_controller` commands Cartesian poses through libfranka’s pose interface, so the contact interaction is mediated by the robot’s internal impedance law rather than a user-tuned stiffness.

Contact force in this thesis is measured through the Franka FCI’s own external-wrench estimate; a second, independent channel using an array of Tekscan FlexiForce resistive force sensors was provisioned, but a repeatable calibration could not be established in time and the cross-check is carried as future work (Chapter 7). FlexiForce is a thin-film piezoresistive sensor marketed explicitly for grip-force measurement in

robotic and surgical applications [36], and has prior academic precedent for exactly this use: Sharan and Onwubolu [37] instrumented a two-finger pick-and-place gripper with a FlexiForce sensor to measure grip force directly against known payload weight. An alternative tactile-sensing approach not used in this thesis but relevant for future work is vision-based tactile sensing: Yuan et al.’s GelSight [38] sensor measures the deformation of a soft elastomer contact surface via an internal camera, recovering high-resolution contact geometry from which force and slip can be inferred, in contrast to FlexiForce’s direct (but lower-resolution) force read-out.

Finally, the trial procedure’s touch-off step (descending the gripper under position control until a chosen contact criterion is met, before switching to the pressed-depth grasp) follows a long lineage of contact-detection strategies in manipulation. Whitney’s classical analysis of quasi-static assembly [39] derives the geometric and force-equilibrium conditions under which measured contact force can be used to detect and guide compliant insertion, the analytical basis for the Remote Center Compliance device and, more generally, for using force as a contact-detection signal rather than only as an outcome measure. De Luca et al. [40] showed that contact can be detected from proprioceptive joint measurements alone, without a dedicated force/torque sensor, via a generalised-momentum residual, a capability this thesis’s FCI-based touch-off implicitly relies on when it uses the controller’s own external-wrench estimate rather than an external load cell.

2.5 The Franka FR3 Platform

The Franka Research 3 (FR3) is a 7-degree-of-freedom, torque-controlled collaborative arm with seven integrated joint torque sensors, a 3 kg payload, and position repeatability under ± 0.1 mm, accessed in this thesis through the Franka Control Interface (FCI) [41, 42], a real-time control layer providing torque, position, and velocity access in both joint and Cartesian space at up to 1 kHz. This thesis controls the arm through `franka_ros2` [43], the official ROS 2 integration built on `libfranka` and `ros2_control`.

The FR3’s predecessor, the Franka Emika Panda, is documented by its manufacturer as a widely adopted reference platform for robotics research and education [44], and a more recent survey by the same author quantifies its continued adoption as a de

facto standard platform across manipulation, control, and human-robot-interaction research [45]. That widespread adoption has produced independent characterisation studies directly relevant to this thesis’s instrumentation choices. Most pertinently, Petrea et al. [46] empirically evaluated the accuracy of the Panda’s built-in filtered external-force estimate against ground-truth load-cell measurements during controlled contact. This is the same FCI external-wrench signal this thesis uses as its Contact Force channel, so their accuracy findings directly bound how much that channel can be trusted, and motivate the independent force-sensor cross-check identified as future work. Gaz et al. [47] separately derived a full dynamic (inertial and friction) parameter model of the Panda’s arm via penalty-based optimisation, a model since reused widely enough to be treated as a de facto standard dynamic model of the platform. Beyond force sensing specifically, the Franka Panda/FR3’s broad adoption as a manipulation-research testbed is illustrated by work such as Chi et al.’s [48] diffusion-policy visuomotor control framework, evaluated on contact-rich real-world manipulation tasks using a Franka Panda as one of its physical platforms. This adoption is directly relevant to the control mode chosen here: learned visuomotor policies of this kind emit streams of end-effector pose targets, the same command interface through which this thesis’s trial procedure drives the arm (Section 3.4). The benchmark is therefore conducted through the interface a learned or simulation-trained policy would use in deployment, with sim-to-real transfer as the motivating trajectory: a compliant finger that tolerates centimetre-scale press-depth error and degree-scale tilt error widens the placement-error budget available to a policy trained in simulation, and the depth-tilt maps produced in this thesis quantify that budget as a function of gripper design.

2.6 Position and Summary

This thesis sits at the intersection of three threads: cloth-handling system design, finray-effect gripper characterisation, and force-controlled manipulation on collaborative robots. Prior work has independently established each thread: cloth manipulation research has largely optimised for task-level success (folding, singulating, dressing) under vision-driven or learned control; compliant-gripper research has characterised finray and jamming designs primarily through force/holding-capacity metrics

on rigid or quasi-rigid test objects; and impedance-control theory and its Franka FCI implementation provide the tools to regulate and measure contact force but have not, in the reviewed literature, been used to map how that force varies over a continuous depth–tilt parameter space for a fragile, non-rigid workpiece. The contribution of this thesis is a unified experimental treatment of all three on a single platform with consistent instrumentation, enabling direct quantitative comparison between rigid and compliant gripper response under deliberately perturbed conditions.

3

Experimental Setup and Methodology

3.1 Hardware Configuration

3.1.1 Robot Manipulator

A Franka Emika FR3 7-DOF cobot served as the manipulator throughout this work. The arm was operated through the Franka Control Interface (FCI) using `libfranka` 0.18.1 and the ROS 2 Humble control stack. Cartesian pose tracking was performed by a custom `fr3_pose_controller` developed and maintained by the McGill Applied Dynamics Lab. The controller accepts a sequence of (pose, twist) waypoints with associated time stamps and interpolates between them using a quintic polynomial; for the in-place re-orientation and the final precision descent phases, a cosine half-cycle S-curve interpolator (the “sine descent”) is used to ensure zero start- and end-state velocities.

3.1.2 Gripper Variants Tested

Table 3.1: Six gripper variants benchmarked, indexed by blade thickness t and baseline seat-mount tilt α (Figure 3.1). All finray variants share the same 3D-printed inclined-seat mount; only α changes between the $+2.5^\circ$ baseline and the $+5.0^\circ$ variant. The parallel-jaw variants do not use an inclined seat and have no defined t or α .

#	Variant	Type	Finger	t	α
1	parallel_jaw_stock	Rigid	Metal jaws (Franka Hand)	—	—
2	parallel_jaw_TPU	Rigid + pad	2.0 mm TPU 90A pad facing	—	—
3	finray_18	Compliant	TPU 90A print	18 mm	$+2.5^\circ$
4	finray_18_model2	Compliant	Second print of #3	18 mm	$+2.5^\circ$
5	finray_26	Compliant	TPU 90A print	26 mm	$+2.5^\circ$
6	finray_26_5deg	Compliant	TPU 90A print, 5.0° seat	26 mm	$+5.0^\circ$

The finray fingers and the parallel-jaw compliant pad were 3D-printed from Bambu Lab TPU 90A (Shore-A hardness 90) on a Bambu Lab X1 Carbon printer using the vendor-default TPU 90A profile in Bambu Studio (0.4 mm nozzle, default layer height, default infill, PEI-textured build plate).

Baseline inclined-seat mount. We define α as the baseline seat-mount tilt angle, annotated in the front view of Figure 3.1: the fixed rotation about the gripper’s local y axis (the gripping direction) introduced by the 3D-printed seat adapter that carries every finray finger. The figure shows the mounting stack (flange $\rightarrow$ seat $\rightarrow$ finger) and the α and t design-parameter convention. All finray variants were mounted on the Franka Hand via this inclined seat; the seat was present in *every* finray variant tested in this thesis unless otherwise noted, and is the physical realisation of α . Variants whose name contains `_5deg` used an $\alpha = +5.0^\circ$ seat in place of the default $\alpha = +2.5^\circ$ seat; all other finray variants (`finray_18`, `finray_18_model2`, `finray_26`) used the $\alpha = +2.5^\circ$ baseline; the parallel-jaw variants have no seat and no defined α . The $\alpha = +2.5^\circ$ baseline was chosen so that the two opposing finray fingers converge inward toward each other rather than splaying apart. This closes the residual slit that would otherwise remain between the finger tips when the gripper is fully closed, and ensures firm contact along the fingers’ full length during a successful grasp.

Consequence for reported “0°” cells. Cells reported in Chapter 5 as 0° tilt mean the robot arm is perpendicular to the surface ($\theta_x = \theta_y = 0$; the runtime tilt angles are illustrated in Figure 3.3). Runtime tilt values ($\pm 2.5^\circ$ about x or y) tilt the arm away from this perpendicular pose.

3.1.3 Cloth Specimen

3.1.4 Test Surface

[TODO: insert photograph of the test bench in situ and short caption describing the table material, flatness, and cloth-edge orientation relative to the gripper y axis.]

3.2 Reference Frame and Tool-Centre Calibration

3.2.1 Frame Convention

Three frames are relevant: the world (robot base) frame $\mathcal{W}$, the gripper flange frame $\mathcal{F}$, and the gripper tip frame $\mathcal{T}$. The Franka URDF defines $\mathcal{F}$ at the mechanical hand mounting flange; the Franka API `robot.end_effector_pose` reports the pose of $\mathcal{F}$ in $\mathcal{W}$. The tip frame $\mathcal{T}$ is located along the local $+z$ axis of $\mathcal{F}$ by a per-variant offset z_{TCP} , which depends on the physical gripper (finger length, mounting fixtures, pads).

Throughout, $\mathbf{p} \in \mathbb{R}^3$ (boldface) denotes a position vector expressed in $\mathcal{W}$ unless otherwise noted, and $R_x(\phi)$, $R_y(\phi)$ denote the elementary rotation matrices for a right-handed rotation by angle ϕ about the fixed world x or y axis, so that a general gripper orientation is a composition $R \in SO(3)$ of these primitives. For a gripper at orientation R , the world-frame relationship between the tip and flange positions is

$$\mathbf{p}_{\mathcal{T}} = \mathbf{p}_{\mathcal{F}} + R \begin{bmatrix} 0 \\ 0 \\ z_{\text{TCP}} \end{bmatrix}. \quad (3.1)$$

For the rigid-down-pointing nominal orientation $R_0 = R_x(180^\circ)$, the tip lies directly below the flange and $\mathbf{p}_{\mathcal{T}} = \mathbf{p}_{\mathcal{F}} - z_{\text{TCP}}\hat{\mathbf{z}}$.

The gripper-frame convention adopted here is: the gripping axis (along which the fingers close) is $\mathcal{F}$'s local y axis; the gripper is mirror-symmetric about the xz plane; consequently, rotation of the gripper about the local y axis is rotation about the geometrically asymmetric axis, while rotation about x is rotation about the symmetric axis. Because the nominal orientation $R_0 = R_x(180^\circ)$ is itself a rotation about the world x axis, the world x and y axis *lines* coincide with the flange's local symmetric and gripping axis lines at R_0 (up to sign). The runtime tilt angles θ_x, θ_y introduced in Section 3.4 (Figure 3.3) are defined as rotations about these fixed world axes, and because the axis lines coincide with the gripper-local axes at R_0 , the same symmetric-axis / gripping-axis language applies without ambiguity to θ_x and θ_y respectively.

3.2.2 TCP Offset Calibration

The z_{TCP} value differs by variant because different fingers attach to the same flange. To calibrate, each variant was driven to a common probe position (x_0, y_0) on a known rigid table and a force-feedback touch-off was performed: the gripper descended slowly in z until the external force in z exceeded a -3 N threshold (debounced for three consecutive 50 Hz samples). The flange z at the latch moment was recorded as $z_{\text{TCP}}^{\text{measured}}$.

Taking the rigid `parallel_jaw_stock` variant as the URDF reference ($z_{\text{TCP}} \equiv 0$ by definition), the per-variant TCP offset is

$$z_{\text{TCP}}^{(v)} = z_{\text{TCP}}^{(v),\text{measured}} - z_{\text{TCP}}^{\text{stock},\text{measured}}. \quad (3.2)$$

Table 3.2 summarises the calibrated values used throughout this work.

Table 3.2: Calibrated per-variant TCP offsets from differential touch-off against the stock parallel jaw.

Variant	$z_{\text{TCP}}^{\text{measured}}$ (mm)	$z_{\text{TCP}}^{(v)}$ (mm)
<code>parallel_jaw_stock</code>	37.871	0.000
<code>parallel_jaw_TPU</code>	39.916	2.044
<code>finray_18</code>	131.842	93.971
<code>finray_26</code>	132.782	94.911
<code>finray_26_5deg</code>	130.094	92.223

The reproducibility of the touch-off was measured as the across-tap standard deviation $\sigma_{z,\text{tcp}}$ in mm (not to be confused with the per-cell force standard deviation σ_{F_z} of Section 3.7). For `finray_26_5deg`, $\sigma_{z,\text{tcp}} = 183\text{ }\mu\text{m}$ owing to a wobble component introduced by the inclined seat; the other four variants had $\sigma_{z,\text{tcp}} \leq 102\text{ }\mu\text{m}$. The second print (`finray_18_model2`) shares the finger geometry and seat of `finray_18`; its surface reference and TCP offset were re-established by the same touch-off procedure when it was mounted, rather than reusing the first print’s calibration.

3.3 Force Measurement

External forces $\mathbf{F}_{\text{ext}}$ at the flange were estimated by the Franka internal model from joint torques, available via the `libfranka` RobotState API at 1 kHz and subsampled to 30 Hz for session logging. A free-air bias capture was performed at the start of every session: the gripper was held at a known pose 30–80 mm above the cloth surface and 30 samples were averaged over 1 s to provide F_z^{bias} . This bias was recorded in each session JSON but *not* subtracted from the raw force time series at acquisition time; downstream analyses subtract it where appropriate.

3.4 Tilt-Aware Pose Targeting

When the gripper is tilted, the flange position required to place the tip at a fixed probe location $(x_0, y_0, z_{\text{surf}})$ shifts laterally. From equation (3.1), the required flange position is

$$\mathbf{p}_{\mathcal{F}}^* = \mathbf{p}_{\mathcal{T}}^{\text{target}} - R \begin{bmatrix} 0 \\ 0 \\ z_{\text{TCP}} \end{bmatrix}. \quad (3.3)$$

This relation is implemented in the trial controller `run_trial.py`: the target orientation $R = R_{\text{tilt}} \circ R_0$ is constructed at session start by composing the requested tilt rotation $R_{\text{tilt}} = R_y(\theta_y)R_x(\theta_x)$ (extrinsic, world frame) onto the base orientation $R_0 = R_x(180^\circ)$, and the flange position target is updated accordingly. The orientation is then held constant through all subsequent descent and grasp phases, so the tip XY remains at (x_0, y_0) throughout the trial. Figure 3.3 illustrates the two runtime tilt angles: θ_x rotates the gripper about the symmetric axis (front view, both fingers move together), while θ_y rotates it about the gripping axis (side view, the mirror-

symmetric fingers overlap in projection). Both are shown here as a general rotation from the world z axis (“Ground”); the runtime sweep used in this study is $\theta_x, \theta_y \in \{-2.5^\circ, 0^\circ, +2.5^\circ\}$. These arm rotation angles are a separate quantity from the seat angle α of Section 3.1.2: α is a fixed property of the gripper hardware (how the fingers meet the flange), while θ_x, θ_y describe how the arm is oriented relative to the surface. The two are reported independently throughout this thesis and are never summed.

For a $+2.5^\circ$ tilt about x on the `finray_26` variant ($z_{\text{TCP}} = 94.911\text{ mm}$), the predicted flange offset is

$$\Delta y_{\mathcal{F}} = -z_{\text{TCP}} \sin(2.5^\circ) \approx -4.14\text{ mm}. \quad (3.4)$$

Chapter 4 reports the empirical lateral offsets measured under load and the residuals from this geometric prediction, which we attribute to finger compliance under press.

3.5 Data Pipeline

3.5.1 Per-Trial Recording

Each trial directory `trial_NNN/` contains:

- `config.yaml` — snapshot of the variant config at trial time;
- `run_config.json` — depth, tilt, ground-truth source, speeds, gripper mode;
- `pose_wrench.csv` — high-rate (30 Hz) actual + commanded pose, RPY, wrench, gripper width per row, tagged with the current trial stage.

3.5.2 Per-Session Recording

Each session directory `depth_*/` contains in addition:

- `depth_summary.json` — header metadata + array of per-trial verdicts (`success`, `fail_reason`); operator notes;
- `ground_truth_used.json` — snapshot of the GT JSON loaded;
- `session_log_*.npz` — numpy archive of the 30 Hz background sample stream, with stage labels and tap indices.

3.5.3 Verdict Recording and Outcome Classification

Trial outcomes were classified into three categories rather than a binary success/failure. After each trial the operator was prompted to distinguish *success* (the gripper held the cloth throughout the lift stage) from *no-grasp* (the trial executed to completion and the stationary force sample was valid, but the cloth slipped from the gripper during or after the lift). Trials in which the FCI control daemon aborted the motion mid-trial were recorded as *failed trials* and excluded from all aggregate analyses; the two most common triggers were:

- the Franka safety reflex firing on excessive predicted joint force (`auto_failed_excess_force`) which typically occurred when the parallel-jaw variants were commanded to depths beyond their narrow force-limited working range;
- the Franka safety reflex firing on acceleration discontinuity between two consecutive Cartesian trajectory segments, which was observed intermittently during the chunked touch-off descent and mitigated by raising the trajectory-emission floor of the touch-off descent trajectories.

Both trigger classes left the trial’s `success` field as `null` in the per-trial JSON; no-grasp trials, by contrast, carried `success = false` with a valid force record and were retained for engagement-force analysis.

3.6 Trial Procedure

A single trial proceeded as follows:

0. Cloth reset: the operator flattened the cloth on the test surface, smoothing out any wrinkles or folds left by the previous trial, so that every trial began from the flat-on-surface configuration the benchmark targets;
1. The gripper opened (Franka Hand or finray release width);
2. Bulk descent: a smooth multi-waypoint trajectory carried the gripper from the lifted pose down to a fine clearance (+5 mm above surface), filtered to exclude any upward leg (preventing direction reversals at zero-twist interior waypoints);

3. Settle: a 2 s pause at the fine-clearance pose;
4. Final approach: a cosine-S-curve descent carried the gripper from the fine clearance to the press target $z_{\text{surf}} - d$, where d is the trial press depth (Figure 3.4);
5. Grasp: a force-controlled close (Franka grasp action at 80 N or 20 N nominal, with ± 100 mm ϵ to disable the position-acceptance check that would otherwise cause auto-release on a sub-spec final width);
6. Settle: a 1 s pause;
7. Stationary sample: 30 samples of $\mathbf{F}_{\text{ext}}$ were captured over 1 s during the `grasp_stationary_s` stage;
8. The gripper lifted to $\sim +80$ mm above surface for visual inspection;
9. The operator was prompted for a verdict;
10. The gripper opened and the rig was reset for the next trial.

3.7 Statistical Treatment

3.7.1 Calibrated Depth Reference

All aggregate analyses are reported on a *calibrated* depth axis

$$d_{\text{cal}} = d - d_{\text{first success}}, \quad (3.5)$$

where $d_{\text{first success}}$ is, per trial series (each variant–tilt-condition pair, within a calibration epoch), the shallowest cell in which at least two-thirds of the valid completed trials succeed (the majority-success criterion); each series is zeroed at its own first majority-success cell. The reference is thus a directly measured physical event — the onset of reliable grasping — rather than the force-touch-off zero, which inherits cross-session bias drift and per-print geometry offsets (Section 3.2.2). Aggregate statistics and figures use only $d_{\text{cal}} \in [0, 10]$ mm; cells outside this window remain in the record but are excluded from aggregates. This convention was adopted as Bill 0007 of the

project’s governance record (enacted 2026-08-09; revised to per-series zeroing and a majority-success cell criterion on 2026-08-10), on the evidence that first-success alignment collapses the two prints of the 18mm finger onto a common curve.

3.7.2 Per-Cell Statistics

For each (variant, depth, tilt) cell, the contact force is reported as

$$\bar{F}_z = \frac{1}{N} \sum_{i=1}^N F_z^{(i)}, \quad \sigma_{F_z} = \sqrt{\frac{1}{N} \sum (F_z^{(i)} - \bar{F}_z)^2}, \quad (3.6)$$

where N is the total number of stationary samples pooled across all clean trials in the cell. Trials with null verdict (FCI crash) are excluded. Trial-to-trial variance is the dominant source of dispersion; the within-trial sample variance (~ 0.1 N) is small compared to it.




Figure 3.1: Both grippers in isometric, front, and side orthographic views, mounted on the Franka Hand. Top: a finray variant; the front view (centre) annotates the baseline seat-mount tilt α at each finger root, and the side view (right) annotates the finger blade thickness t . Bottom: the rigid parallel-jaw gripper with the optional TPU tip indicated. Scale bars in cm.






Figure 3.2: The single-layer plain-woven cotton specimen used throughout this thesis, 30×22.5 cm, shown against a ruler for scale. The same physical specimen was reused across all trials, reset to a flat, wrinkle-free configuration before each trial (Section 3.6, step 0).



(a) θ_x : rotation about the symmetric axis (front view)



(b) θ_y : rotation about the gripping axis (side view)

Figure 3.3: Runtime tilt angles θ_x (left) and θ_y (right), each shown as a rotation of the whole gripper body away from the world z axis (labelled) relative to the Ground line; 0° is the arm perpendicular to the surface.




Figure 3.4: Physical effect of the press depth d on a finray variant. At first contact ($d = 0$, left) the blades are straight and only the tips touch the surface; pressed to $d > 0$ (right), the compliant blades buckle inward and wrap the surface, increasing the contact area. This wrap-around engagement underlies the depth-gated tilt robustness discussed in Chapter 6.

4

Experimental Design and Procedure

The empirical campaign was organised into two sequential phases against a single experimental apparatus. Phase 1 established the contact-force versus press-depth relationship for each gripper variant at nominal zero-tilt orientation, and Phase 2 perturbed the same measurement with $\pm 2.5^\circ$ tilts about the two orthogonal Cartesian axes. Together the two phases define a three-dimensional (depth, tilt-axis, tilt-sign) cell grid whose per-cell measurements constitute the data corpus reported in Chapter 5.

4.1 Study Design

The gripper roster comprised six variants organised into two families:

- **Finray family** (four variants): `finray_18`, `finray_18_model2` (second print of the same CAD), `finray_26` (larger blade), and `finray_26_5deg` (`finray_26` on a 5° -inclined seat).

- **Parallel-jaw family** (two variants): `parallel_jaw_stock` (bare rigid jaws) and `parallel_jaw_TPU` (rigid jaws with a 2 mm TPU 90A compliant pad).

The same fabric specimen (single-layer cotton) rested on the same acrylic table surface throughout the campaign. The Franka FR3 arm held the gripper in a fixed base orientation $R_0 = R_x(180^\circ)$ for Phase 1, and applied signed rotations $R_\theta = R_x(180^\circ) R_y(\theta_y) R_x(\theta_x)$ for Phase 2, with $(\theta_x, \theta_y) \in \{(\pm 2.5^\circ, 0), (0, \pm 2.5^\circ)\}$ (the runtime arm tilts of Chapter 3, Figure 3.3, distinct from the seat angle α of Figure 3.1).

Table 4.1 summarises the test conditions common to both phases.

Table 4.1: Test conditions of the experimental campaign. Symbols are defined in the Notation section; θ_x, θ_y are illustrated in Figure 3.3.

Variable	Definition	Values tested
Press depth d	Fingertip driven below the detected surface	rigid: -1 to $+4$ mm at 0.25 – 0.5 mm steps finray: -2 to $+16$ mm at 0.5 – 1.0 mm steps
Tilt θ_x	Arm leaned about the symmetric axis	$-2.5^\circ, 0^\circ, +2.5^\circ$
Tilt θ_y	Arm leaned along the gripping direction	$-2.5^\circ, 0^\circ, +2.5^\circ$
Surface reference	Force touch-off at 3 N threshold	repeatable to ± 0.1 mm
Repeats	Trials per (variant, depth, tilt) cell	3

4.2 Depth Grid

The depth axis is defined relative to the per-session detected cloth surface via the `ground_truth_finder` touch-off procedure described in Chapter 3. The depth grid was chosen adaptively per gripper family:

- For the rigid parallel jaws, depths were spaced at 0.25 – 0.5 mm intervals over -1 to $+4$ mm, close to the narrow force-limited working range.
- For the finrays, depths were spaced at 0.5 – 1.0 mm intervals over -2 to $+16$ mm, covering the four regimes (boundary, knee, valley, recovery) identified in Chapter 5.

Per (variant, depth, tilt) cell, three trials were collected. Boundary depths were routinely retested when within-cell standard deviation exceeded 1 N, up to a maximum of two additional trials per (depth, tilt) combination.

4.3 Tilt-Aware Pose Geometry: Verification

Before reporting force results, the implementation of the tilt-aware tool-centre-point (TCP) tracking equation was verified empirically. Table 4.2 shows the measured flange XY shift during the stationary grasp sample versus the geometric prediction for `finray_26` at four representative (depth, tilt) cells.

Table 4.2: Measured vs. predicted flange-XY offset during the stationary grasp sample at four cells. Predicted offset uses $z_{\text{TCP}} = 94.911$ mm.

Cell	Tilt cmd	Predicted $\Delta \mathbf{p}_{\mathcal{F}}$ (mm)	Measured $\Delta \mathbf{p}_{\mathcal{F}}$ (mm)	Residual	Mean F_z
−1.0 mm	+2.5° y	(+4.14, +0.00)	(+4.35, −0.03)	0.21 mm	−8.00 N
−1.0 mm	−2.5° y	(−4.14, +0.00)	(−4.02, −0.04)	0.13 mm	−3.39 N
+10.0 mm	+2.5° y	(+4.14, +0.00)	(+5.35, −0.28)	1.25 mm	−36.26 N
+10.0 mm	−2.5° y	(−4.14, +0.00)	(−2.93, −0.29)	1.21 mm	−35.95 N

At boundary depth, the residual is ≤ 0.21 mm—essentially the lateral positioning noise of the arm under light load. At safe depth, both positive and negative tilts exhibit a +1.2 mm common-mode offset in the same $+x$ direction, regardless of tilt sign. We interpret this common-mode offset as evidence of finger compliance under load: the ~ 36 N axial press deflects the soft finray blade laterally by approximately 1.2 mm in a structurally preferred direction. The differential offset (between $+y$ and $-y$ tilts) remains correctly predicted by the geometric model in both depth conditions.

This verification confirms that the control-frame implementation of the tilt-aware TCP tracker is correct, and any asymmetries observed in the contact force reflect physical gripper behaviour rather than control-system errors.

4.4 Rotation Tracking Verification

A natural concern is whether the Franka arm faithfully realises the commanded tilt orientation. The 30 Hz session log records actual and commanded RPY at every

sample, enabling a direct measurement of orientation tracking error

$$\delta(t) = \|\text{RPY}_{\text{actual}}(t) - \text{RPY}_{\text{cmd}}(t)\|_2. \quad (4.1)$$

Across five representative sessions, the mean steady-state δ was 0.05° with a worst-case 0.13° . Transient excursions up to $\sim 5^\circ$ occur exclusively within the `orientation_lock` stage at session start, before any descent or grasp, and are filtered out of trial-level analyses.

A 2.5° commanded tilt is therefore $50\times$ the noise floor of the arm’s rotation tracking; the observed force differences cannot be attributed to arm rotation error.

4.5 Trial Protocol and Outcome Classification

Each trial followed the fixed stage sequence detailed in Section 3.6: set the commanded pose above the target, descend to the pre-contact height, contact and press to the target depth, grasp with the gripper’s default close command, lift by a fixed clearance, and record the operator’s held/not-held verdict. A stationary-sample window of 1 s (30 samples) during the grasped-and-pressed stage provided the force measurement.

Trial outcomes were classified into three categories:

Success: the trial executed to completion, contact force was sampled cleanly during the stationary window, and the cloth remained gripped after the lift stage.

No-grasp: the trial executed to completion and contact force was sampled cleanly during the stationary window, but the cloth slipped from the fingers during or after the lift stage. The force measurement remains valid as an *engagement-force* characterisation.

Failed trial: the trial did not execute to completion. The controller either aborted with an excess-force safety event or the operator interrupted before the stationary window completed. The measurement is invalid and the trial is excluded from analyses.

The rationale for this three-way taxonomy—rather than a binary success/failure classification—is discussed in Section 6.5.

4.6 Experiment 3: Inclined-Surface Demonstration (Pilot)

Experiments 1 and 2 characterise the contact phase on a flat, levelled surface, with each variant calibrated against that surface. Experiment 3 tests the deployment consequence those measurements predict: when the surface geometry departs from what the calibration assumed, does the compliant finger keep grasping where the rigid jaw fails, with no change to calibration, policy, or sensing?

The planned protocol mirrors Experiments 1 and 2 in every respect except the surface and the outcome metric. The acrylic test surface is mounted on a wedge (SURFACE_RISER) whose CAD model specifies a top-surface incline of $\beta = 4.15^\circ$ (measured from the model geometry: the top face normal is 4.1467° from vertical), with the slope aligned with the gripping direction; the as-built fixture will be verified against this nominal value before the session. Each variant is then calibrated *blind* by the standard force touch-off of Section 3.2.2, exactly as if the surface were flat and level: no tilt compensation is applied and the incline is treated as unknown to the calibration. Each trial presses to the calibrated touch-off depth plus 0.5 mm, the shallowest depth at which every variant grasped successfully on the flat surface, and the outcome is recorded as a binary grasp verdict under the outcome taxonomy of Section 4.5. Contact force is logged but is not a reported metric for this experiment.

The pilot covers two variants, `finray_26` and `parallel_jaw_stock`, at three attempts each. From the flat-surface working ranges of Chapter 5, the registered prediction is that the finray variant grasps successfully at all three attempts, while the rigid jaw fails or is marginal: a 4.15° surface-relative misalignment exceeds the largest runtime tilt at which the rigid jaws retained a usable window (2.5°), and that window had already narrowed to between 0.5 and 2.5 mm under that tilt.

[TODO: Experiment 3 is a placeholder: the protocol above is enacted but the pilot has not yet been run. Insert the outcome grid (variant $\times$ attempt, success/no-grasp/abort), a photograph of the wedge fixture with its measured incline angle, and a short results paragraph once the session is complete. If the pilot is not run before submission, delete this section and retain the inclined-surface demonstration as future work in Chapter 7.]

5

Results

This chapter reports the empirical findings across all six gripper variants (`finray_18`, `finray_18_model2`, `finray_26`, `finray_26_5deg`, `parallel_jaw_stock`, `parallel_jaw_TPU`) organised around the primary measurement (contact force F_z) at swept command depths and $\pm 2.5^\circ$ tilts about the x and y axes. All values were F_z -bias-corrected using the per-session Franka F_{ext} zero-load reading and pooled across trials with matching command depths and tilts unless otherwise noted. Interpretation is deferred to Chapter 6.

5.1 Canonical Force-Depth Curves at 0°

Under zero tilt, all four finray variants exhibited a stereotyped *bell-plus-valley* force-depth response with four regions: (i) a boundary regime where contact force rose steeply with press depth, (ii) a knee peak around +1.5 to +3 mm past the detected surface, (iii) a valley trough near +6 to +8 mm, and (iv) a recovery regime deeper than +10 mm. Table 5.1 summarises the key features.

The four features lay within ± 1 mm of the same depth locations across all four

Table 5.1: Key features of the 0° force-depth curve for each finray variant. Peak $|F_z|$ is the maximum observed contact force in the knee region; valley $|F_z|$ is the local minimum between the knee peak and the deep-regime recovery.

Variant	Knee peak $ F_z $	Peak depth	Valley $ F_z $	Valley depth
finray_18	35.4 N	+1.5 mm	31.9 N	+6 mm
finray_18_model2	31.2 N	+3.0 mm	25.2 N	+8 mm
finray_26	~ 45 N	+2.0 mm	~ 41 N	+6 mm
finray_26_5deg	~ 50 N	+2.5 mm	~ 43 N	+6 mm

variants. The absolute force magnitudes scaled approximately with the frontal contact area of the finger.

Normalising each variant’s curve by its own peak force and averaging over its five tested orientations makes this shared shape directly visible (Figure 5.1): all four finray variants collapse onto a single rise–peak–valley–recovery curve over the full $[0, 10]$ mm calibrated window. Table 5.2 quantifies the pairwise agreement between variants as a normalised mean-squared error (NMSE) and, in decibels, $10 \log_{10}(\text{NMSE})$; the mean pairwise NMSE is 0.58% (≈ -22 dB), with the largest disagreement between any two variants at 0.80% (-21.0 dB).

Table 5.2: Pairwise shape agreement between finray variants over the full $[0, 10]$ mm calibrated window: NMSE (top), NMSE in decibels (middle), and RMSE as a percentage of peak (bottom, in brackets).

	finray_18	finray_18_model2	finray_26	finray_26_5deg
		0.41%	0.69%	0.80%
finray_18	—	(-23.9 dB)	(-21.6 dB)	(-21.0 dB)
		[5.5%]	[7.2%]	[7.8%]
			0.13%	0.70%
finray_18_model2		—	(-28.8 dB)	(-21.6 dB)
			[3.1%]	[7.1%]
				0.76%
finray_26			—	(-21.2 dB)
				[7.4%]
finray_26_5deg				—




Figure 5.1: Peak-normalised $|F_z|$ versus calibrated depth d_{cal} , averaged over all five tested orientations per variant (shaded band: across-orientation spread). All four finray variants collapse onto one shared response shape.

5.2 Bipolar Tilt Bias on the Finray Family

A $\pm 2.5^\circ$ tilt produced a signed bias $b(d, a) = |F_z^{+2.5^\circ}(d, a)| - |F_z^{-2.5^\circ}(d, a)|$ whose depth trajectory itself followed a bipolar shape: a strong signed peak at boundary depths, a zero-crossing near the knee, a sign-reversed plateau through the valley, and depth-gated attenuation toward zero at deep press. The clearest example was the y-axis bias on `finray_26_5deg`, which had full coverage across 14 sign-flip pairs from -2 to $+16$ mm:

- Boundary peak: $b = +5.40 \pm 0.12$ N at $d = -1.5$ mm.
- Zero crossing: between $d = -0.5$ mm ($b = +2.00$ N) and $d = +1.5$ mm ($b = -1.35$ N).
- Knee plateau: $b = -1.57 \pm 0.21$ N at $d = +2.0$ mm, persisting to $d = +3.0$ mm.
- Depth-gated recovery: $|b| < 0.9$ N at all depths beyond $d = +4$ mm.

The x-axis bias on `finray_26_5deg` had a similar bipolar shape but with opposite sign at the boundary ($b_x = -4.17 \pm 0.32$ N at $d = -1.5$ mm).

5.3 Cross-Axis Sign Reversal at the Boundary

At boundary depths, the x-axis and y-axis biases had *opposite signs* on `finray_26_5deg` (Table 5.3).

Table 5.3: Bias on `finray_26_5deg` at boundary depths. x-axis and y-axis biases had opposite signs. Values are $|F_z^{+2.5^\circ}| - |F_z^{-2.5^\circ}|$ in newtons.

Depth	x-bias	y-bias	Sign product
-1.5 mm	-4.17 ± 0.32	$+5.40 \pm 0.12$	—
-1.0 mm	-3.94 ± 0.18	$+4.27 \pm 0.07$	—
-0.5 mm	-3.51 ± 0.39	$+2.00 \pm 0.24$	—

The same sign-reversal pattern was present on `finray_18` at $d = -1.0$ mm (x -bias = -2.44 ± 0.17 N; y -bias = $+4.49 \pm 0.47$ N), demonstrating that the effect was not specific to the 5° mount condition.

5.4 Mount-Tilt Alignment: 5° Mount as a Depth Shift

After a two-parameter best-fit alignment (single depth shift plus single force offset), the 5° -mounted variant collapsed onto the flat-mounted variant with residual RMSE below 2 N on both tilt axes independently (Table 5.4).

Table 5.4: Best-fit alignment of `finray_26_5deg` onto `finray_26`. The two axes yielded independent estimates of the depth shift within 0.03 mm of each other.

Axis	Depth shift	Force offset	Alignment RMSE
x	-1.14 mm	+2.7 N	1.75 N
y	-1.12 mm	+3.1 N	1.58 N
<i>mean</i>	-1.13 mm	+2.9 N	1.66 N

5.5 Tilt-Aligned Response Comparison

Plotting each finray variant’s five tilt-condition curves together on the calibrated depth axis (Figure 5.2) shows how tightly the orientations track one another within a variant. The mean spread across the five orientations, averaged over the calibrated window, was 1.3–2.4 N depending on variant — small relative to the 30–51 N knee peaks reported in Table 5.1, and consistent with the depth-gated tilt robustness reported in Section 5.2: the curves separate near the boundary and converge in the plateau.




Figure 5.2: Contact force versus calibrated depth for each finray variant, all five tested orientations overlaid (0° , $x \pm 2.5^\circ$, $y \pm 2.5^\circ$). Within each panel the five curves nearly coincide.

5.6 Print-to-Print Manufacturing Variability

Figure 5.3 compares the two prints directly on the calibrated depth axis. Two independent prints of the same `finray_18` CAD geometry, each zeroed at its own first majority-success cell (Section 3.7.1), differed by a +4.6 N mean force offset (second print stiffer) with RMSE 5.4 N over the 10 mm window; removing the constant offset left a residual RMSE of 2.9 N. This ≈ 5 N offset was adopted as the canonical same-session print-to-print variability, approximately 15% of the knee peak force.

Cross-session drift on the *same* print, measured by comparing `finray_18` data collected on separate days approximately three weeks apart, was approximately 7–10 N at knee and knee-adjacent depths. Cross-session drift on the same specimen was therefore larger than same-session print-to-print variability.




Figure 5.3: Contact force versus calibrated depth for two independently printed copies of `finray_18`, zero tilt, successful cells only. Error bars: one standard deviation across repeats.

The two prints also differed qualitatively in x-axis bias sign: `finray_18` exhibited a bipolar x-bias (negative at boundary, positive at knee), whereas `finray_18_model12` exhibited a monotonically negative x-bias across the entire measured range.

5.7 Trial Outcome Distribution

Trial outcomes were classified into three categories: *success* (cloth held after lift), *no-grasp* (trial completed with valid F_z measurement but cloth slipped at lift), and *failed trial* (controller aborted or trial did not complete a stationary window). The failed-trial category was excluded from force-curve analyses; no-grasp trials contributed valid force measurements. No-grasp fractions at each variant’s boundary depth are reported in Table 5.5.

Figure 5.4 shows the same information as a continuous function of calibrated depth at zero tilt: grasp success rate climbs from 0% at the no-contact boundary

Table 5.5: No-grasp fraction at each variant’s boundary depth. A no-grasp trial yields a cleanly measurable contact force but fails to hold the cloth on lift.

Variant	Boundary depth	No-grasp fraction
finray_18	−2.0 mm	3/3
finray_18_model2	−1.0 mm	3/3
finray_26	−0.5 mm	~ 40%
finray_26_5deg	−2.0 mm	6/6
parallel_jaw_stock	+1.0 mm	4/6
parallel_jaw_TPU	+0.5 mm	varies with tilt

to 100% by the first majority-success cell and holds at 100% for every finray variant across the remainder of the $[0, 10]$ mm window.




Figure 5.4: Grasp success rate versus calibrated depth, zero tilt, binned in 0.5 mm increments below 3 mm and pooled above. $\times$: 0% success.

5.8 Working-Range and Failure Comparison

Figure 5.5 summarises the zero-tilt working range of every variant, and Figure 5.6 extends the comparison to all five tested orientations. On the calibrated depth axis (Section 3.7.1), the parallel-jaw variants exhibited closed successful-grasp windows of 1.5 mm (parallel_jaw_stock) and 2.5 mm (parallel_jaw_TPU) at 0° , shrinking to 0.5–2.5 mm under tilt, with two stock-jaw tilt conditions producing no successful

grasp at any depth. All four finray variants spanned the full ≥ 10 mm analysis window under every orientation tested; their ranges were right-censored at the window end, and the underlying record extends beyond it.




Figure 5.5: Successful-grasp working range on the calibrated depth axis, zero tilt. Hatched bars are right-censored at the 10 mm analysis-window end (Section 3.7.1).

The parallel jaws auto-failed on the Franka excess-force safety monitor at depths only 1–1.5 mm beyond their successful-grasp window; the finrays never triggered the safety monitor within the tested -2 to $+16$ mm range on any variant.




Figure 5.6: Successful-grasp working range on the calibrated depth axis, all five tested orientations per variant. ×: no successful attempt at any depth. Hatched bars are right-censored at the analysis-window end.

6

Discussion

The results reported in Chapter 5 are interpreted in thematic groups mirroring the results sections, with an emphasis on the physical mechanisms that produce the observed force behaviour and the implications for cloth-handling system design.

6.1 Bell-Plus-Valley Curve Shape

The bell-plus-valley force-depth shape observed on all four finray variants (Section 5.1) is consistent with the finray finger operating in three sequential contact regimes as press depth increases. In the boundary regime, only the distal tip of the finger contacts the cloth; force rises roughly linearly with penetration because the contact area grows in proportion to displacement. At the knee, the full ridged section of the finger is engaged and the finger resists compression through its cross-web structure, producing the peak force. Past the peak, the ribs partially buckle and the finger folds back on itself; contact area increases but effective stiffness drops, producing the valley. At recovery depths, the finger has fully wrapped and further penetration is resisted by structural compression through the finger base, producing the second rise.

The finding that the four features (boundary knee, peak, valley, recovery) occur at the same command depths (± 1 mm) across all four finray variants indicates that this three-regime response is a geometric property of the finray shape itself and not sensitive to finger thickness (18 vs 26 mm) or to a 5° mount tilt. Force magnitudes scale with contact area, as expected, but the depth locations of the transitions do not.

6.2 Cross-Axis Sign Reversal: Mechanism

The finray finger has *mirror symmetry* about the x-axis (the mirror plane between the two opposing fingers) but not about the y-axis (the gripping direction, which breaks left-right equivalence by construction). The observed y-axis bias at boundary depths is therefore consistent with the intrinsic finger geometry: a $+2.5^\circ$ rotation about the gripping axis loads the two fingers differently and produces a signed force difference.

The x-axis bias is more surprising. If the finger and cloth were both symmetric about the mirror plane, a $\pm 2.5^\circ$ rotation about the mirror axis should produce zero bias by symmetry. That the x-bias is nonzero *and of opposite sign to the y-bias* indicates the symmetry-breaking is not on the finger side. The most parsimonious explanation is that the cloth specimen has a directional weave-warp axis that is not aligned with the mirror plane of the gripper. A $+2.5^\circ$ x-rotation of the gripper then produces a systematic cloth-load asymmetry that couples through the whole-hand contact into a signed force difference. The finding that this x-bias is reproducible across two independently mounted variants (`finray_18` and `finray_26_5deg`) is consistent with a shared external source (cloth or table) rather than a per-print effect.

The prediction from this mechanism is that rotating the cloth specimen 90° in the lab frame should flip the sign of the x-bias while leaving the y-bias approximately unchanged (because the intrinsic finger asymmetry does not depend on the cloth orientation). This is the single most important follow-up experiment to isolate the cloth contribution.

6.3 Mount-Tilt Equivalence: Geometric Interpretation

The observation that a 5° pre-tilt of the finray mount is equivalent to a -1.13 mm depth shift, with axis-independent shift within 0.03 mm on the two axes (Table 5.4), admits a simple geometric interpretation. A 5° tilt of a rigid finger of length L produces a projected vertical offset at the contact tip of $L \sin(5^\circ) = 0.087L$. Setting this equal to the observed 1.13 mm shift gives $L_{\text{eff}} \approx 13$ mm. Because the finray finger is physically ~ 50 mm long, this L_{eff} corresponds to a contact-engagement point roughly 25% up the length of the finger from its distal tip, consistent with the load being concentrated on the proximal portion of the finray’s curved contact surface rather than at the tip.

The axis-independence of the shift (0.02 mm apart on the two axes) is a strong constraint: it rules out any mechanism in which the pre-tilt produces a different effective load-bearing length on the two axes. The 5° mount therefore acts as a strictly rigid translation of the operating point in the depth dimension, independent of the tilt axis applied at runtime. This makes the 5° mount a useful design lever: the finray’s operating envelope can be shifted a chosen 1.13 mm in depth without altering the underlying force-depth sensitivity, at the cost of the ~ 3 N vertical force offset.

6.4 Manufacturing Variability vs Session Drift

The finding that same-day print-to-print variability (≈ 4.6 N mean, RMSE 5.4 N) is *smaller* than same-print cross-session drift (7 – 10 N) has two consequences worth stating explicitly.

First, from a design-decision perspective, any cross-variant force-magnitude claim with an effect size below the ≈ 5 N print-to-print offset is unlikely to reproduce across different physical instances of the same CAD, and should be regarded as tentative until replicated across prints before being carried into product decisions.

Second, from an experimental-methodology perspective, session drift is the dominant noise source in cloth-grasping benchmarks of this kind. Any measurement campaign that pools data across days must either report cross-day drift as a separate variance component or restrict its headline claims to same-session pairs. In this

thesis, all bias values are reported for same-session pairs unless explicitly labelled otherwise; cross-day comparisons are used only for shape-consistency observations, not for headline effect sizes.

The qualitative x-axis bias sign flip between `finray_18` and `finray_18_model2` (bipolar vs monotonically negative) is a puzzle that this thesis cannot resolve with the current data. Two prints from the same slicer, same infill, same material, and same printer produced measurably different tilt-response behaviours on the x-axis but similar behaviour on the y-axis. The most likely mechanism is small differences in rib-wall thickness or infill density that manifest as asymmetric buckling behaviour under x-axis load, but verifying this requires post-hoc CT-scanning or a controlled study with more prints. The finding is noted here as an interesting boundary condition on manufacturing-invariance claims for compliant grippers.

6.5 Failure-Mode Taxonomy: Why the Distinction Matters

The trial-outcome taxonomy of success, no-grasp, and failed trial (Section 5.7) is deliberately three-valued rather than binary because the three categories map to physically distinct regimes with distinct control implications.

Treating no-grasp trials as failed trials would systematically delete data from the boundary regime, where cloth engagement is intermittent but force measurement is clean. Boundary-regime force values are essential for characterising the low end of the working range, so their exclusion would truncate the observed force-depth curves and mask the boundary-tilt biases that are one of the headline findings.

The opposite error, treating failed trials as no-grasps, would pollute the force averages with spurious near-zero readings from aborted controllers, inflating apparent variance and eroding the signal-to-noise of the bias estimates. The three-way taxonomy is the minimum classification that separates “the trial produced valid measurement” from “the grasp was successful” from “the trial was invalid.”

For deployment implications, the no-grasp regime is precisely where a force-feedback grasp controller could intervene productively: the contact force is being measured and the depth is known, but the lift is failing. A closed-loop controller

monitoring $|F_z|$ during lift could detect the impending slip and either deepen the press or abort before dropping the cloth outside the pick zone. Distinguishing no-grasp from failed-trial in the data record makes this control strategy visible in the offline analysis.

6.6 Deployment-Envelope Implications for Cloth-Handling Systems

The parallel-jaw versus finray comparison (Section 5.8) puts at least a $4\times$ – $7\times$ factor between the two classes in tolerance to depth command noise — a lower bound, the finray windows being right-censored at the analysis-window end. This factor is large enough that any cloth-handling system whose depth uncertainty exceeds 1 mm should default to a finray class gripper unless there is a specific force-precision requirement that the finray cannot meet.

The absence of excess-force safety events on any finray variant across the full -2 to $+16$ mm range indicates that finray-based systems can be deployed open-loop with no depth clipping controller and no safety-limit tuning. Parallel-jaw systems require an active depth governor whose accuracy budget must be at least a factor of two tighter than the successful-grasp window, i.e. sub-half-millimetre command precision, to avoid frequent aborts.

The expanded dataset supports a three-position design envelope — safe-plateau finray, boundary finray with tilt control, and rigid jaw with compliant pad — with one refinement: the compliant-pad variant (`parallel_jaw_TPU`) has narrower successful-grasp coverage than its compliant facing might suggest, and its excess-force susceptibility is closer to `parallel_jaw_stock` than to finray. The compliant-pad design remains a useful hybrid but its robustness benefit is qualitative (better boundary behaviour) rather than transforming the parallel jaw into a finray substitute.

6.7 Limitations

Single cloth specimen and single surface. The reported working ranges and cloth-fixture asymmetries were measured on one cotton cloth specimen on one acrylic table surface. The absolute numbers will shift with cloth thickness, weave direction,

friction, and surface compliance. The specific cross-axis sign reversal attributed to cloth weave in Section 6.2 requires a 90°-rotated cloth control to isolate the cloth contribution from the fixture contribution.

Cross-session drift. The dominant noise source in this dataset is cross-session drift on the same specimen, exceeding print-to-print manufacturing variability. All same-day comparisons reported here are internally valid; cross-day claims are qualified throughout.

Parallel-jaw x-axis coverage. The parallel-jaw variants were not swept across the x-axis tilt sign; the parallel-jaw bias values are restricted to the y-axis. Whether the x-axis bipolar bias observed on the finrays has an analogue on the rigid grippers is untested.

Boundary-depth sample sizes. Boundary-depth cells frequently have $n = 3$ and inter-trial standard deviations of 0.5–1.5 N driven by cloth-fold state variability that is not further characterised. Boundary values are honest measurements with correspondingly wide error bars.

Success rate as a binary rating. The success/no-grasp distinction is operator-assessed after the lift stage; a finer-grained grip-force-during-lift measurement would resolve marginal cases that the current binary rating collapses.

6.8 Code and Data Availability

All source code, raw trial data, ground-truth JSON files, and analysis scripts are published in the project repository at [TODO: public GitHub URL]. The thesis was written against commit [TODO: commit hash].

Conclusion and Future Work

7.1 Answers to the Research Questions

This thesis presented a controlled benchmark of six gripper variants on a single platform (Franka FR3), measuring stationary contact force and grasp outcome as a function of press depth and tilt angle over 845 analysed trials. Returning to the four research questions posed in Chapter 1:

1. *How does contact force vary with press depth for each finger design?* Every finray variant follows a single shared rise–peak–valley–recovery curve (pairwise NMSE below 1% over the full 10 mm calibrated window), whereas the rigid parallel jaws ramp steeply toward the platform’s excess-force limit within roughly 2 mm of first contact.
2. *How much can the gripper be tilted before the grasp becomes unreliable?* The finrays grasp reliably across every orientation tested ($\pm 2.5^\circ$ about both axes, plus a 5° seat mount) over the full ≥ 10 mm window; the rigid jaws never exceed

2.5 mm, and two stock-jaw tilt conditions produced no successful grasp at any depth.

3. *Does tilt tolerance depend on press depth?* Yes — robustness is depth-gated: near the contact boundary, tilt strongly modulates both force and success (including a $2.4\times$ direction-dependent asymmetry about the geometrically asymmetric axis), whereas past the buckling threshold the mean across-orientation force spread shrinks to 1.3–2.4 N.
4. *What is the compliant finger’s response under compression?* Blade buckling: force peaks at buckling onset ($\approx 2\text{--}3\text{ mm}$ past first successful grasp), dips as the blades buckle, then recovers — and the finray geometry (blade thickness, seat angle, print instance) sets the force *scale* (30–51 N peaks) but not the curve *shape*.

The depth-dependent tilt robustness identified here had not been quantitatively characterised in prior literature, to the author’s knowledge, and constitutes the principal empirical contribution.

7.2 Implications

With the robot, control policy, cloth, and environment held identical, a finger swap alone buys at least a $4\times$ wider press-depth tolerance than the best rigid configuration ($\geq 10\text{ mm}$ versus 2.5 mm on the padded jaws, and roughly $7\times$ versus the stock jaws), at roughly half the grip force. Under tilt the advantage widens further, to $20\times$ or more against the stock jaws’ surviving conditions; all of these ratios are lower bounds, because the finray ranges are right-censored at the analysis-window end. For the application domains that motivated this work (laundry and textile handling, hospital bedding, assistive robotics, garment manufacturing, and soft-goods packaging), this means the flat-cloth pick, the step that currently keeps such systems human-started, becomes accessible to a standard arm running an unchanged policy, before any investment in cameras, tactile sensing, or learned control. The accompanying design rule is equally simple: press the compliant finger past its buckling peak ($\approx 2\text{--}3\text{ mm}$ beyond first contact), and the grasp becomes robust to both depth and tilt placement error.

7.3 Future Work

Independent force-sensor cross-check. All force measurements derive from the arm’s joint-torque-based external-wrench estimate. Repeating a subset of the matrix with a dedicated force/torque sensor at the wrist (or an instrumented surface) would bound the estimator’s pose-dependent bias directly.

Task-level evaluation. Downstream task-level evaluation — folding a cloth diagonally from a corner grasp, dropping a centre-grasped cloth onto a curved object, dragging between positions under the same policy — would test whether the contact-phase advantage measured here translates into end-to-end task success, including under learned policies.

Task-level precision comparison. Beyond success rates, a soft versus rigid comparison of task-level *precision* (placement accuracy after fold or drop) would characterise the cost side of compliance.

Dynamic and high-rate behaviour. All trials here are quasi-static. The perturbation behaviour of compliant fingers under fast actuation and high-frequency control — where the compliance that absorbs placement error may instead cause oscillation or delayed settling — is uncharacterised.

Tilt-sensitivity model. An analytical model relating finger curvature, contact compliance, and tilt-induced contact-point migration would connect the present empirical findings to the broader continuum-mechanics literature on soft contact.

Multi-fabric study. All trials used a single cloth specimen. Replicating the matrix on fabrics of varying thickness, weave, and surface treatment would establish the applicability range of the depth–tilt–robustness curves identified here.

Closed-loop control. The boundary-depth regime, where finray directional asymmetry dominates, motivates a tilt-compensating controller using inline force feedback, validated against the open-loop boundary baseline.

Sim2real validation. A simulation of the finray finger in a compliant-contact physics engine could be calibrated against the empirical force–depth–tilt surface from

this thesis, providing a digital twin for the gripper variants.



Detailed Coverage and Supplementary Data

This appendix records the final per-variant coverage of the experimental campaign. The table below was regenerated programmatically from the on-disk trial record (`experiments/cloth_grasp/`) on 2026-08-04, counting only active cells (staged, superseded, and backup cells excluded). A *cell* is one (variant, depth, tilt) session directory; each cell contains between one and several trials.

A.1 Per-Variant Coverage Table

In addition to the conditions tabulated above, `finray_26` carries five exploratory cells at $\theta_x = +5.0^\circ$ (depths -0.5 to $+10$ mm, recorded 2026-06-05) and one zero-tilt boundary-depth cell recorded within the Experiment 2 campaign; these were pilot probes of the larger-tilt regime and are not part of the primary $\pm 2.5^\circ$ analysis.

Across both experiments the active record comprises 945 trials. This figure counts every trial directory on disk, including no-grasp outcomes (retained as valid data per

Table A.1: Experiment 1 (depth sweep at zero commanded tilt): active cell counts, sampled depth range, and total trials per variant.

Variant	Cells	Depth range (mm)	Trials
finray_18	17	$[-2, +14]$	51
finray_18_model2	28	$[-1, +14]$	86
finray_26	18	$[-1, +22]$	50
finray_26_5deg	10	$[-1, +20]$	30
parallel_jaw_TPU	9	$[-1, +4]$	26
parallel_jaw_stock	12	$[-1, +4]$	29
Total	94		272

the outcome taxonomy of Section 3.5.3) and failed trials (excluded from aggregate analyses); the trial counts quoted in Chapter 5 are the subset surviving the per-analysis exclusion rules stated there.

Table A.2: Experiment 2 (tilt sweeps): active cell counts per commanded tilt condition, with sampled depth range in mm in brackets, and total trials per variant. The parallel-jaw variants could only be swept over their narrow force-limited depth window, and their $\theta_x = -2.5^\circ$ condition is covered by mirror symmetry of the jaw about the xz plane rather than by separate cells.

Variant	$\theta_x=+2.5^\circ$	$\theta_x=-2.5^\circ$	$\theta_y=+2.5^\circ$	$\theta_y=-2.5^\circ$	Trials
finray_18	12 $[-1, +12]$	13 $[-2, +12]$	11 $[-2, +12]$	10 $[-2, +12]$	134
finray_18_model2	12 $[-2, +12]$	12 $[-2, +12]$	12 $[-2, +12]$	11 $[-2, +12]$	140
finray_26	14 $[0, +16]$	14 $[0, +16]$	11 $[-2, +16]$	10 $[-2, +16]$	163
finray_26_5deg	15 $[-2, +16]$	17 $[-2, +16]$	14 $[-2, +16]$	14 $[-2, +16]$	170
parallel_jaw_TPU	6 $[0, +4]$	—	5 $[0, +4]$	3 $[0, +3]$	40
parallel_jaw_stock	4 $[+1, +2]$	—	3 $[+1, +2]$	3 $[+1, +2]$	26
Total					673

List of Publications

No portion of this thesis has been published, presented, or submitted for publication at the time of writing. A manuscript covering the depth–tilt benchmark results (Chapters 5 and 6) is in preparation.

Bibliography

- [1] E. Musk, “Optimus folds a shirt.” X (formerly Twitter) post, Jan. 2024. Video, posted Jan. 15, 2024; accessed Aug. 12, 2026. A follow-up post by the author acknowledged the robot was teleoperated and the setup fixed.
- [2] M. Novak, “Elon musk’s latest robot video accidentally gives away the magic trick.” Forbes, Jan. 2024. Published Jan. 15, 2024; accessed Aug. 12, 2026.
- [3] Y. Kita, F. Saito, and N. Kita, “A deformable model driven visual method for handling clothes,” in *2004 IEEE International Conference on Robotics and Automation (ICRA)*, (New Orleans, LA), pp. 3889–3895, 2004.
- [4] J. Maitin-Shepard, M. Cusumano-Towner, J. Lei, and P. Abbeel, “Cloth grasp point detection based on multiple-view geometric cues with application to robotic towel folding,” in *2010 IEEE International Conference on Robotics and Automation (ICRA)*, pp. 2308–2315, 2010.
- [5] S. Miller, J. van den Berg, M. Fritz, T. Darrell, K. Goldberg, and P. Abbeel, “A geometric approach to robotic laundry folding,” *The International Journal of Robotics Research*, vol. 31, no. 2, pp. 249–267, 2012.
- [6] Y. Li, D. Xu, Y. Yue, Y. Wang, S.-F. Chang, E. Grinspun, and P. K. Allen, “Regrasping and unfolding of garments using predictive thin shell modeling,” in *2015 IEEE International Conference on Robotics and Automation (ICRA)*, pp. 1382–1388, 2015.
- [7] K. Saxena and T. Shibata, “Garment recognition and grasping point detection for clothing assistance task using deep learning,” in *2019 IEEE/SICE International Symposium on System Integration (SII)*, pp. 632–637, 2019.

- [8] Y. Wu, W. Yan, T. Kurutach, L. Pinto, and P. Abbeel, “Learning to manipulate deformable objects without demonstrations,” in *Robotics: Science and Systems (RSS) XVI*, 2020.
- [9] S. Tirumala, T. Weng, D. Seita, O. Kroemer, Z. Temel, and D. Held, “Learning to singulate layers of cloth using tactile feedback,” in *2022 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2022.
- [10] A. Longhini, Y. Wang, I. Garcia-Camacho, D. Blanco-Mulero, M. Moletta, M. Welle, G. Alenyà, H. Yin, Z. Erickson, D. Held, J. Borràs, and D. Kragic, “Unfolding the literature: A review of robotic cloth manipulation,” *Annual Review of Control, Robotics, and Autonomous Systems*, vol. 8, pp. 295–322, 2025.
- [11] A. Ramisa, G. Alenyà, F. Moreno-Noguer, and C. Torras, “Using depth and appearance features for informed robot grasping of highly wrinkled clothes,” in *2012 IEEE International Conference on Robotics and Automation (ICRA)*, pp. 1703–1708, 2012.
- [12] J. Qian, T. Weng, L. Zhang, B. Okorn, and D. Held, “Cloth region segmentation for robust grasp selection,” in *2020 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pp. 9553–9560, 2020.
- [13] H. Ha and S. Song, “Flingbot: The unreasonable effectiveness of dynamic manipulation for cloth unfolding,” in *Proceedings of the 5th Conference on Robot Learning (CoRL)*, vol. 164 of *Proceedings of Machine Learning Research*, 2021.
- [14] G. J. Monkman, “Robot grippers for use with fibrous materials,” *The International Journal of Robotics Research*, vol. 14, no. 2, pp. 144–151, 1995.
- [15] J. Guo, J. Leng, and J. Rossiter, “Electroadhesion technologies for robotics: A comprehensive review,” *IEEE Transactions on Robotics*, vol. 36, no. 2, pp. 313–327, 2020.
- [16] B. Ozcelik and F. Erzincanli, “A non-contact end-effector for the handling of garments,” *Robotica*, vol. 20, no. 4, pp. 447–450, 2002.

- [17] H. Jiang, X. Jiang, Y. Zhou, H. Chen, P. Li, and Y. Liu, “Variable-friction-based in-hand manipulation of fabrics applied to unfolding operations,” *IEEE Robotics and Automation Letters*, vol. 7, no. 4, pp. 12259–12266, 2022.
- [18] L. Li, D. Crosby, M. Shuttleworth, O. F. Argin, A. S. Chen, G. Herrmann, R. Kay, and A. Weightman, “A comparative analysis and scoping review of soft-rigid and industrial parallel rigid grippers,” *Advanced Intelligent Systems*, vol. 7, p. 2400503, 2025.
- [19] A. B. Clark, L. Cramphorn-Neal, M. Rachowiecki, and A. Gregg-Smith, “Household clothing set and benchmarks for characterising end-effector cloth manipulation,” in *2023 IEEE International Conference on Robotics and Automation (ICRA)*, 2023.
- [20] L. Kniese, “Load carrying element with flexible outer skin,” 2000. European Patent Application EP1040999A2.
- [21] Festo AG & Co. KG, “Adaptive gripper fingers DHAS (technical data sheet),” 2018.
- [22] X. Shan and L. Birglen, “Modeling and analysis of soft robotic fingers using the fin ray effect,” *The International Journal of Robotics Research*, vol. 39, no. 14, pp. 1686–1705, 2020.
- [23] W. Crooks, G. Vukasin, M. O’Sullivan, W. Messner, and C. Rogers, “Fin ray® effect inspired soft robotic gripper: From the robosoft grand challenge toward optimization,” *Frontiers in Robotics and AI*, vol. 3, p. 70, 2016.
- [24] K. Elgeneidy, P. Lightbody, S. Pearson, and G. Neumann, “Characterising 3d-printed soft fin ray robotic fingers with layer jamming capability for delicate grasping,” in *2019 2nd IEEE International Conference on Soft Robotics (RoboSoft)*, (Seoul, South Korea), pp. 143–148, 2019.
- [25] L. Lang, R. Antunes, M. Lima de Aguiar, N. Pereira, T. Assis Dutra, Y. Lu, and P. D. Gaspar, “Soft robotics for advanced handling of delicate fruit products,”

- in *Proceedings of the 22nd International Conference on Informatics in Control, Automation and Robotics (ICINCO)*, vol. 2, pp. 569–576, 2025.
- [26] A. M. Dollar and R. D. Howe, “The highly adaptive SDM hand: Design and performance evaluation,” *The International Journal of Robotics Research*, vol. 29, no. 5, pp. 585–597, 2010.
 - [27] C. B. Teeple, J. Werfel, and R. J. Wood, “Multi-dimensional compliance of soft grippers enables gentle interaction with thin, flexible objects,” in *2022 IEEE International Conference on Robotics and Automation (ICRA)*, pp. 728–734, 2022.
 - [28] E. Brown, N. Rodenberg, J. Amend, A. Mozeika, E. Steltz, M. R. Zakin, H. Lipson, and H. M. Jaeger, “Universal robotic gripper based on the jamming of granular material,” *Proceedings of the National Academy of Sciences*, vol. 107, no. 44, pp. 18809–18814, 2010.
 - [29] J. R. Amend, E. Brown, N. Rodenberg, H. M. Jaeger, and H. Lipson, “A positive pressure universal gripper based on the jamming of granular material,” *IEEE Transactions on Robotics*, vol. 28, no. 2, pp. 341–350, 2012.
 - [30] F. Ilievski, A. D. Mazzeo, R. F. Shepherd, X. Chen, and G. M. Whitesides, “Soft robotics for chemists,” *Angewandte Chemie International Edition*, vol. 50, no. 8, pp. 1890–1895, 2011.
 - [31] Soft Robotics Inc., “Soft robotics expands mGrip modular gripping system with new features,” 2021. Press release, PR Newswire.
 - [32] N. Hogan, “Impedance control: An approach to manipulation: Part I—theory,” *ASME Journal of Dynamic Systems, Measurement, and Control*, vol. 107, no. 1, pp. 1–7, 1985.
 - [33] N. Hogan, “Impedance control: An approach to manipulation: Part II—implementation,” *ASME Journal of Dynamic Systems, Measurement, and Control*, vol. 107, no. 1, pp. 8–16, 1985.

- [34] N. Hogan, “Impedance control: An approach to manipulation: Part III—applications,” *ASME Journal of Dynamic Systems, Measurement, and Control*, vol. 107, no. 1, pp. 17–24, 1985.
- [35] Franka Robotics GmbH, “libfranka: C++ client library for the franka control interface,” 2026. Official software documentation.
- [36] Tekscan, Inc., “Grip force measurement with FlexiForce,” 2026. Official vendor application documentation.
- [37] R. V. Sharan and G. C. Onwubolu, “A low-cost force measurement solution applicable for robotic grippers,” in *Innovations and Advances in Computing, Informatics, Systems Sciences, Networking and Engineering*, vol. 313 of *Lecture Notes in Electrical Engineering*, pp. 95–98, Springer International Publishing, 2015.
- [38] W. Yuan, S. Dong, and E. H. Adelson, “GelSight: High-resolution robot tactile sensors for estimating geometry and force,” *Sensors*, vol. 17, no. 12, p. 2762, 2017.
- [39] D. E. Whitney, “Quasi-static assembly of compliantly supported rigid parts,” *ASME Journal of Dynamic Systems, Measurement, and Control*, vol. 104, no. 1, pp. 65–77, 1982.
- [40] A. De Luca, A. Albu-Schäffer, S. Haddadin, and G. Hirzinger, “Collision detection and safe reaction with the DLR-III lightweight manipulator arm,” in *2006 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, (Beijing, China), pp. 1623–1630, 2006.
- [41] Franka Robotics GmbH, “Franka research 3 (product specifications),” 2026. Official product page.
- [42] Franka Robotics GmbH, “Franka control interface (FCI) documentation,” 2026. Official documentation site.
- [43] Franka Robotics GmbH, “franka_ros2: ROS 2 integration for franka research robots,” 2026. Official GitHub repository, Apache 2.0 license.

- [44] S. Haddadin, S. Parusel, L. Johannismeier, S. Golz, S. Gabl, F. Walch, M. Sabaghian, C. Jähne, L. Hausperger, and S. Haddadin, “The franka emika robot: A reference platform for robotics research and education,” *IEEE Robotics & Automation Magazine*, vol. 29, no. 2, pp. 46–64, 2022.
- [45] S. Haddadin, “The franka emika robot: A standard platform in robotics research,” *IEEE Robotics and Automation Magazine*, vol. 31, no. 4, pp. 136–148, 2024.
- [46] R. A. B. Petrea, M. Bertoni, and R. Oboe, “On the interaction force sensing accuracy of Franka Emika Panda robot,” in *IECON 2021 – 47th Annual Conference of the IEEE Industrial Electronics Society*, pp. 1–6, 2021.
- [47] C. Gaz, M. Cagnetti, A. Oliva, P. Robuffo Giordano, and A. De Luca, “Dynamic identification of the Franka Emika Panda robot with retrieval of feasible parameters using penalty-based optimization,” *IEEE Robotics and Automation Letters*, vol. 4, no. 4, pp. 4147–4154, 2019.
- [48] C. Chi, Z. Xu, S. Feng, E. Cousineau, Y. Du, B. Burchfiel, R. Tedrake, and S. Song, “Diffusion policy: Visuomotor policy learning via action diffusion,” in *Robotics: Science and Systems (RSS)*, 2023.